

CEE6501 — Lecture 4.3

Matrix Bandwidth and Extending the DSM to 3D

Learning Objectives

By the end of this lecture, you will be able to:

- Explain why global stiffness matrices for trusses are typically **sparse**
- Define **half-bandwidth** and relate it to storage and computational cost
- Show how **node numbering / DOF ordering** changes bandwidth without changing physics
- Demonstrate (with timing) why exploiting structure can accelerate linear solves
- Extend the Direct Stiffness Method (DSM) from 2D to **3D trusses**
- Write the **3D truss element** stiffness in local form, the 3D transformation matrix, and the assembled global element stiffness
- State the **support constraints** needed to prevent rigid body motion in 3D

Agenda

1. Sparsity in the global stiffness matrix
2. Half-bandwidth and why node numbering matters
3. A timing demo: same size, different bandwidth
4. DSM in 3D: DOFs, rotations, and supports
5. 3D truss element matrices: \mathbf{k}_{local} , \mathbf{T} , \mathbf{k}_{global}

Part 1 — Sparsity and Bandwidth

In large structural systems, performance depends not only on how many nonzeros we have, but on where they appear in the matrix.

Sparsity (local physical connectivity)

In truss and frame models, each node connects to only a small number of neighboring elements. As a result, each equilibrium equation involves only a few degrees of freedom.

This locality leads to a **sparse** global stiffness matrix \mathbf{K} : most entries are exactly zero.

- Sparsity enables memory-efficient storage
- It also enables specialized **sparse solvers** that avoid unnecessary operations

Sparsity is a direct consequence of the *physics and topology* of the structure.

Bandwidth (connectivity meets indexing)

Sparsity describes *how many* nonzero entries exist. **Bandwidth** describes *how far from the diagonal* those nonzeros extend.

For a symmetric matrix \mathbf{K} , the **half-bandwidth** is defined as

$$NHB = \max\{|i - j| : K_{ij} \neq 0\}.$$

A smaller half-bandwidth means that nonzero entries are clustered closer to the diagonal.

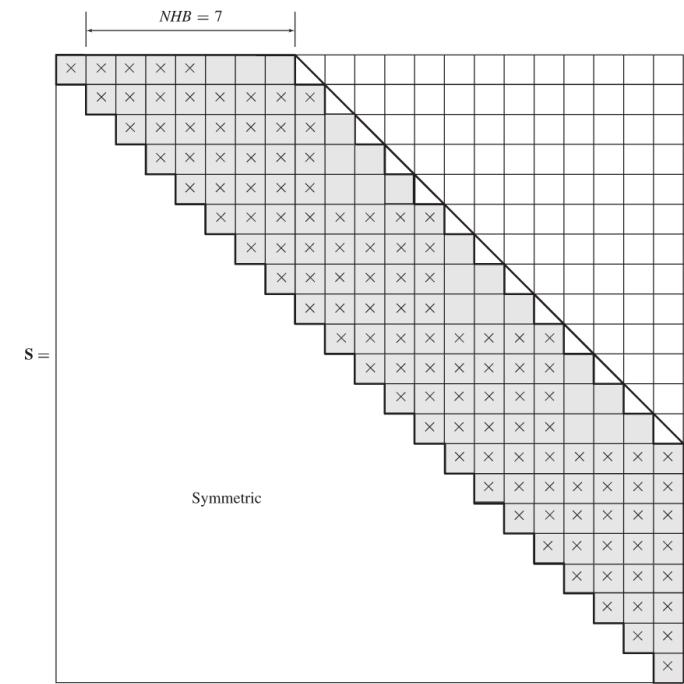
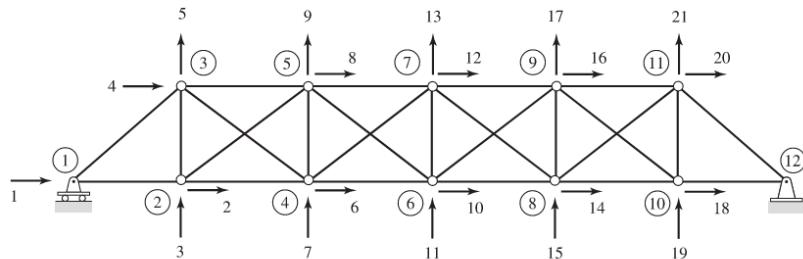
Why node numbering matters

The physical structure is unchanged, but the **algebraic representation** of \mathbf{K} depends on the ordering of degrees of freedom.

- Poor numbering places strongly coupled DOFs far apart → **large bandwidth**
- Good numbering keeps coupled DOFs close → **small bandwidth**

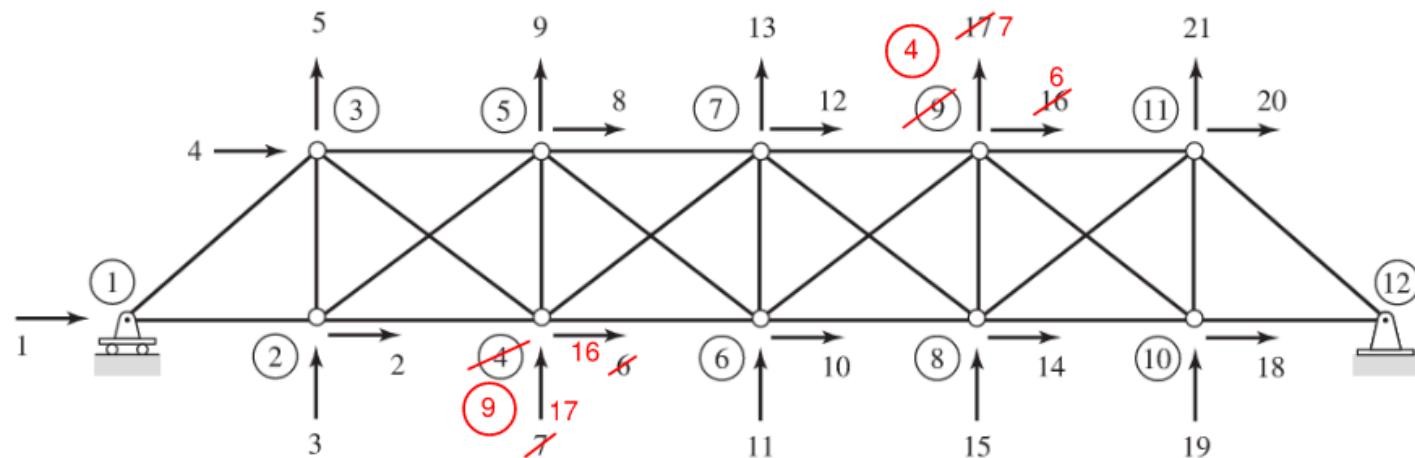
Bandwidth affects the cost of matrix factorization, even though it has no effect on the underlying physics.

Example 1: A well-numbered system

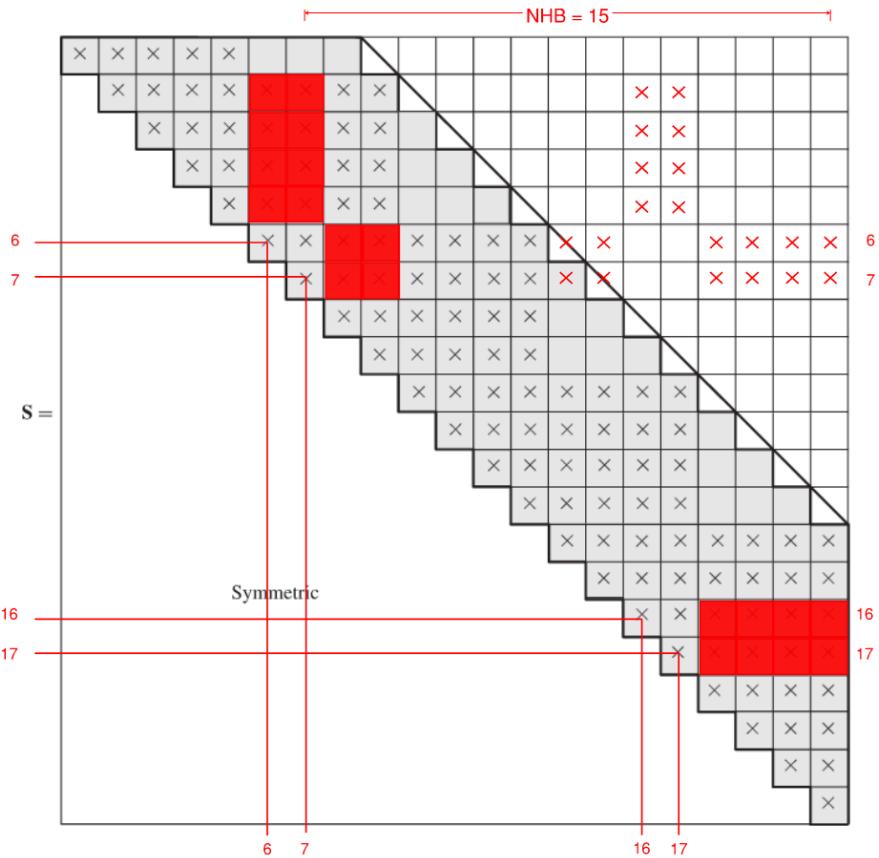
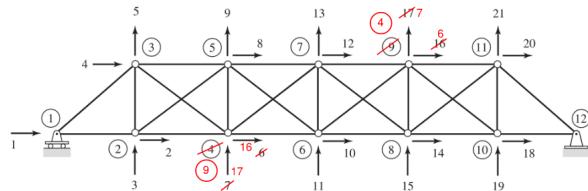


Example 2: A poorly-numbered system

Swap Node 6 and 9, number DOFs accordingly



Half-Bandwidth Increases from 7 to 15



Timing Demo — Same system, different bandwidth

Goal: compare solve times for the **same underlying SPD system** under two different degree-of-freedom orderings.

We construct matrices with:

- identical size n
- identical physical connectivity and sparsity pattern
- but very different **effective half-bandwidths**

This isolates the computational impact of bandwidth and ordering.

Solver types used in the timing demo

We solve the **same linear system** using three strategies that make different assumptions about matrix structure.

1) Dense solve (NumPy)

```
np.linalg.solve(K, B)
```

- Treats \mathbf{K} as fully dense (LU-type)
- Ignores sparsity, bandwidth, and symmetry

2) Sparse LU (no reordering)

```
splu(K, perm_c_spec="NATURAL")
```

- Sparse storage, original DOF ordering
- Strongly affected by bandwidth and fill-in

3) Sparse LU (with reordering)

```
splu(K, perm_c_spec="COLAMD")
```

- Applies fill-reducing reordering
- Reduces bandwidth and factorization cost

Details of sparse factorization are beyond this lecture (see Kassimali §9.9).

Switch to demo code

What this demo shows

- The two systems represent the **same physics**, differing only in DOF ordering
- Dense solvers ignore sparsity and bandwidth entirely (unaffected)
- Larger bandwidth typically leads to higher factorization cost for sparse solvers
- Sparse solvers are sensitive to ordering, but can reduce its impact via reordering

Takeaway: DOF ordering changes the algebraic structure of \mathbf{K} , which can strongly influence computational cost even though the physics is unchanged.

Part 2 — Extending the DSM to 3D Trusses

Same DSM workflow, but with 3 translational DOFs per node and a 3D orientation.

Degrees of freedom in 3D

For a 3D truss node:

$$\mathbf{d}_i = [u_{xi} \quad u_{yi} \quad u_{zi}]^T, \quad \mathbf{d}_j = [u_{xj} \quad u_{yj} \quad u_{zj}]^T$$

Element displacement vector in global coordinates:

$$\mathbf{d}_e = [u_{xi} \quad u_{yi} \quad u_{zi} \quad u_{xj} \quad u_{yj} \quad u_{zj}]^T \equiv \text{DOFs } [1, 2, 3, 4, 5, 6]$$

3D direction cosines

Let the element connect nodes i and j with coordinates (x_i, y_i, z_i) and (x_j, y_j, z_j) .

$$L = \sqrt{(x_j - x_i)^2 + (y_j - y_i)^2 + (z_j - z_i)^2}$$

$$l = \frac{x_j - x_i}{L}, \quad m = \frac{y_j - y_i}{L}, \quad n = \frac{z_j - z_i}{L}$$

These are the direction cosines of the element axis in global coordinates.

3D Truss Element Matrices

We will write:

- the **local** element stiffness \mathbf{k}_{local} (6×6)
- the **transformation** matrix \mathbf{T} (6×6)
- the **global** element stiffness \mathbf{k}_{global} (6×6)

Local stiffness matrix (6×6)

A truss element carries **axial force only** and therefore has stiffness only in the direction of its axis.

In a local coordinate system where the element axis is the local x' direction, only the axial DOFs are active.

Then the 6×6 local stiffness is:

$$\mathbf{k}_{local} = \frac{EA}{L} \begin{bmatrix} 1 & 0 & 0 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Only axial coupling appears; transverse DOFs have zero stiffness in a truss model.

Transformation Matrix: Local axes and direction cosines

A 3D truss element carries **axial force only**, but we still need a consistent mapping between global and local translation components.

Let the element run from node i to node j with coordinates (x_i, y_i, z_i) and (x_j, y_j, z_j) .

Define the element differences and length:

$$\Delta x = x_j - x_i, \quad \Delta y = y_j - y_i, \quad \Delta z = z_j - z_i,$$

$$L = \sqrt{\Delta x^2 + \Delta y^2 + \Delta z^2}.$$

The **axial direction cosines** (local x' axis) are:

$$l = \frac{\Delta x}{L}, \quad m = \frac{\Delta y}{L}, \quad n = \frac{\Delta z}{L}.$$

Why we need *two more* directions in 3D

In 2D, specifying the element axis automatically defines the perpendicular direction. In 3D, the axis (l, m, n) alone does **not** uniquely define a coordinate system: the element can still rotate about its own axis.

So we complete an orthonormal local basis:

- $\hat{\mathbf{e}}_{x'}$ along the member (given by (l, m, n))
- $\hat{\mathbf{e}}_{y'}$ transverse
- $\hat{\mathbf{e}}_{z'}$ transverse

These transverse directions do **not** add stiffness in a truss (still axial-only), but they make the transformation well-defined.

Constructing $\hat{\mathbf{e}}_{y'}$ and $\hat{\mathbf{e}}_{z'}$

Pick a reference vector \mathbf{a} that is **not** parallel to the bar axis. A simple robust rule:

$$\mathbf{a} = \begin{cases} \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}, & |n| < 0.9 \\ \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}, & |n| \geq 0.9 \quad (\text{bar nearly vertical}) \end{cases}$$

Define the local unit vectors:

$$\hat{\mathbf{e}}_{x'} = \begin{bmatrix} l \\ m \\ n \end{bmatrix}$$

$$\hat{\mathbf{e}}_{y'} = \frac{\mathbf{a} \times \hat{\mathbf{e}}_{x'}}{\|\mathbf{a} \times \hat{\mathbf{e}}_{x'}\|}$$

$$\hat{\mathbf{e}}_{z'} = \hat{\mathbf{e}}_{x'} \times \hat{\mathbf{e}}_{y'}$$

Write the transverse direction cosines as:

$$\hat{\mathbf{e}}_{y'} = \begin{bmatrix} l_y \\ m_y \\ n_y \end{bmatrix} \quad \hat{\mathbf{e}}_{z'} = \begin{bmatrix} l_z \\ m_z \\ n_z \end{bmatrix}$$

Meaning: l_y is the x -component of the local y' axis, m_y is the y -component, etc.

3×3 rotation matrix \mathbf{R} (direction cosines)

Collect the three local unit vectors (written in global components) into:

$$\mathbf{R} = \begin{bmatrix} l & m & n \\ l_y & m_y & n_y \\ l_z & m_z & n_z \end{bmatrix}$$

By construction, the basis is orthonormal, so $\mathbf{R}\mathbf{R}^T = \mathbf{I}$.

6×6 transformation matrix \mathbf{T}

Use global element DOFs ordered as:

$$\mathbf{d}_{global} = [u_{xi}, u_{yi}, u_{zi}, u_{xj}, u_{yj}, u_{zj}]^T$$

Use local element DOFs ordered as:

$$\mathbf{d}_{local} = [u_{x'i}, u_{y'i}, u_{z'i}, u_{x'j}, u_{y'j}, u_{z'j}]^T$$

Then:

$$\mathbf{d}_{local} = \mathbf{T} \mathbf{d}_{global}$$

with the block-diagonal transformation:

$$\mathbf{T} = \begin{bmatrix} \mathbf{R} & \mathbf{0} \\ \mathbf{0} & \mathbf{R} \end{bmatrix}$$

Explicitly (in direction cosines):

$$\mathbf{T} = \begin{bmatrix} l & m & n & 0 & 0 & 0 \\ l_y & m_y & n_y & 0 & 0 & 0 \\ l_z & m_z & n_z & 0 & 0 & 0 \\ 0 & 0 & 0 & l & m & n \\ 0 & 0 & 0 & l_y & m_y & n_y \\ 0 & 0 & 0 & l_z & m_z & n_z \end{bmatrix}$$

How \mathbf{T} connects to the 3D truss stiffness

Local (axial-only) stiffness in the element basis:

$$\mathbf{k}_{local} = \frac{EA}{L} \begin{bmatrix} 1 & 0 & 0 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Transformation to global coordinates:

$$\mathbf{k}_{global} = \mathbf{T}^T \mathbf{k}_{local} \mathbf{T}$$

Global element stiffness (6×6)

Compute:

$$\mathbf{k}_{global} = \mathbf{T}^T \mathbf{k}_{local} \mathbf{T}$$

For a 3D truss element, the result can be written directly as:

$$\mathbf{k}_{global} = \frac{EA}{L} \begin{bmatrix} l^2 & lm & ln & -l^2 & -lm & -ln \\ lm & m^2 & mn & -lm & -m^2 & -mn \\ ln & mn & n^2 & -ln & -mn & -n^2 \\ -l^2 & -lm & -ln & l^2 & lm & ln \\ -lm & -m^2 & -mn & lm & m^2 & mn \\ -ln & -mn & -n^2 & ln & mn & n^2 \end{bmatrix}$$

This is the standard 3D truss element stiffness in global coordinates. Even though \mathbf{T} is a full 3D rotation, a truss remains axial-only. That is why the closed-form \mathbf{k}_{global} only contains l, m, n terms.

In [8]:

```
import sympy as sp

# --- symbols ---
EA_L = 1

l, m, n = sp.symbols("l m n", real=True)
ly, my, ny = sp.symbols("l_y m_y n_y", real=True)
```

```

lz, mz, nz = sp.symbols("l_z m_z n_z", real=True)

# --- T matrix (6x6) ---
T = sp.Matrix([
    [l, m, n, 0, 0, 0],
    [ly, my, ny, 0, 0, 0],
    [lz, mz, nz, 0, 0, 0],
    [0, 0, 0, l, m, n],
    [0, 0, 0, ly, my, ny],
    [0, 0, 0, lz, mz, nz],
])
# --- local stiffness (6x6) ---
k_local = EA_L * sp.Matrix([
    [1, 0, 0, -1, 0, 0],
    [0, 0, 0, 0, 0, 0],
    [0, 0, 0, 0, 0, 0],
    [-1, 0, 0, 1, 0, 0],
    [0, 0, 0, 0, 0, 0],
    [0, 0, 0, 0, 0, 0],
])
# --- global stiffness via transformation ---
k_global = sp.simplify(T.T * k_local * T)

# Pretty output (in notebook)
sp.pprint(k_global)

```

$$\left[\begin{array}{cccccc} 2 & & & & & \\ 1 & 1 \cdot m & 1 \cdot n & -1 & 2 & -1 \cdot m & -1 \cdot n \\ & m & m \cdot n & -1 \cdot m & -m & -m \cdot n & \\ 1 \cdot m & & & & & & \\ 1 \cdot n & m \cdot n & n & -1 \cdot n & -m \cdot n & -n & \\ & -m & -m \cdot n & 1 \cdot m & m & m \cdot n & \\ -1 & -1 \cdot m & -1 \cdot n & 1 & 2 & 1 \cdot m & 1 \cdot n \\ & -n & -n \cdot m & 1 \cdot n & m \cdot n & n & \\ -1 \cdot n & -m \cdot n & -n & 1 \cdot n & m \cdot n & n & \end{array} \right]$$

Supports and stability in 3D

A free 3D structure has **3 rigid body translations**.

For a 3D **truss** model (translations only), you must prevent rigid body motion by constraining enough nodal translations.

Typical sufficient constraint patterns include (examples):

- Fix all three translations at one node, and two at a second node, and one at a third node (if geometry allows)
- Or use support conditions consistent with the physical supports (pins/rollers) but ensuring no global drift

If constraints are insufficient, **K** is singular → the model has a mechanism.