

CEE6501 — Lecture 4.3

Matrix Bandwidth and Extending the DSM to 3D

Learning Objectives

By the end of this lecture, you will be able to:

- Explain why global stiffness matrices for trusses are typically **sparse**
- Define **half-bandwidth** and relate it to storage and computational cost
- Show how **node numbering / DOF ordering** changes bandwidth without changing physics
- Demonstrate (with timing) why exploiting structure can accelerate linear solves
- Extend the Direct Stiffness Method (DSM) from 2D to **3D trusses**
- Write the **3D truss element** stiffness in local form, the 3D transformation matrix, and the assembled global element stiffness
- State the **support constraints** needed to prevent rigid body motion in 3D

Agenda

1. Sparsity in the global stiffness matrix
2. Half-bandwidth and why node numbering matters
3. A timing demo: same size, different bandwidth
4. DSM in 3D: DOFs, rotations, and supports
5. 3D truss element matrices: \mathbf{k}_{local} , \mathbf{T} , \mathbf{k}_{global}

Part 1 — Sparsity and Bandwidth

In large structural systems, performance depends not only on how many nonzeros we have, but on where they appear in the matrix.

Sparsity (local physical connectivity)

In truss and frame models, each node connects to only a small number of neighboring elements. As a result, each equilibrium equation involves only a few degrees of freedom.

This locality leads to a **sparse** global stiffness matrix \mathbf{K} : most entries are exactly zero.

- Sparsity enables memory-efficient storage
- It also enables specialized **sparse solvers** that avoid unnecessary operations

Sparsity is a direct consequence of the *physics and topology* of the structure.

Bandwidth (connectivity meets indexing)

Sparsity describes *how many* nonzero entries exist. **Bandwidth** describes *how far from the diagonal* those nonzeros extend.

For a symmetric matrix \mathbf{K} , the **half-bandwidth** is defined as

$$NHB = \max\{|i - j| : K_{ij} \neq 0\}.$$

A smaller half-bandwidth means that nonzero entries are clustered closer to the diagonal.

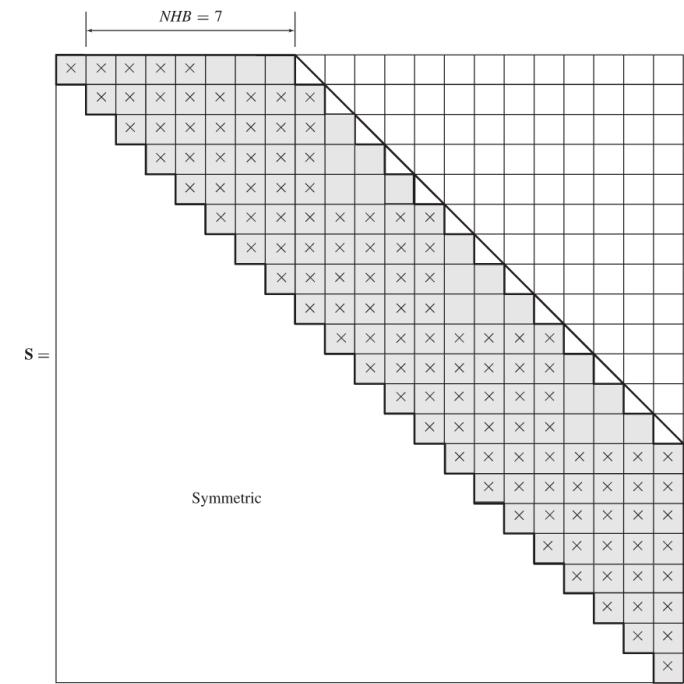
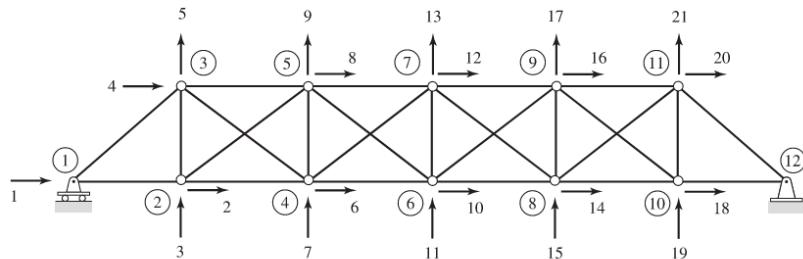
Why node numbering matters

The physical structure is unchanged, but the **algebraic representation** of \mathbf{K} depends on the ordering of degrees of freedom.

- Poor numbering places strongly coupled DOFs far apart → **large bandwidth**
- Good numbering keeps coupled DOFs close → **small bandwidth**

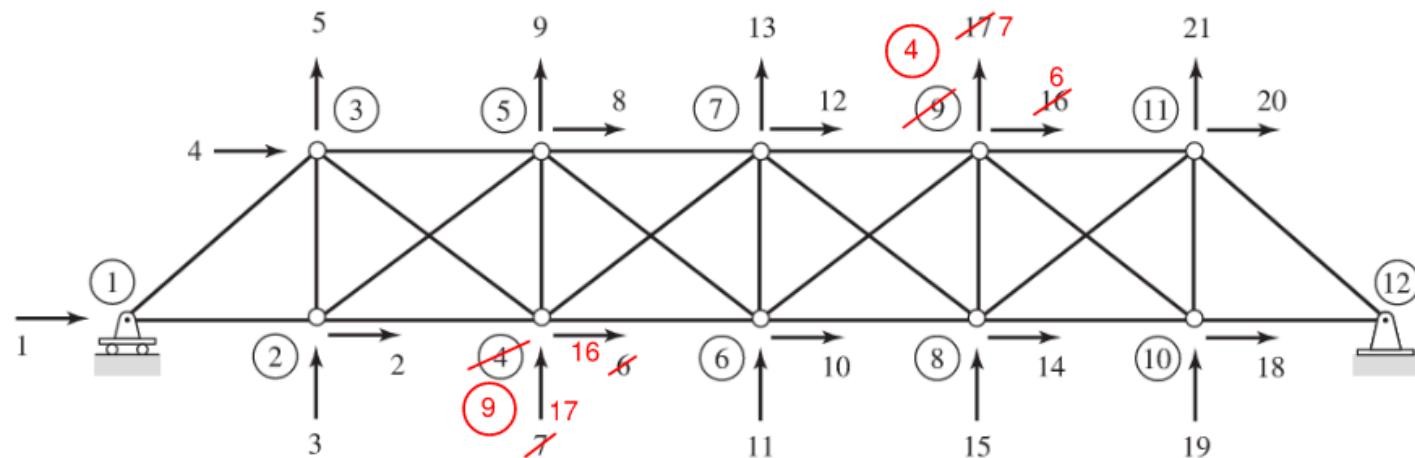
Bandwidth affects the cost of matrix factorization, even though it has no effect on the underlying physics.

Example 1: A well-numbered system





Example 2: A poorly-numbered system

Swap Node 6 and 9, number DOFs accordingly



Half-Bandwidth Increases from 7 to 15



Timing Demo — Same system, different bandwidth

Goal: compare solve times for the **same underlying SPD system** under two different degree-of-freedom orderings.

We construct matrices with:

- identical size n
- identical physical connectivity and sparsity pattern
- but very different **effective half-bandwidths**

This isolates the computational impact of bandwidth and ordering.

Solver types used in the timing demo

We solve the **same linear system** using three strategies that make different assumptions about matrix structure.

1) Dense solve (NumPy)

```
np.linalg.solve(K, B)
```

- Treats \mathbf{K} as fully dense (LU-type)
- Ignores sparsity, bandwidth, and symmetry

2) Sparse LU (no reordering)

```
splu(K, perm_c_spec="NATURAL")
```

- Sparse storage, original DOF ordering
- Strongly affected by bandwidth and fill-in

3) Sparse LU (with reordering)

```
splu(K, perm_c_spec="COLAMD")
```

- Applies fill-reducing reordering
- Reduces bandwidth and factorization cost

Details of sparse factorization are beyond this lecture (see Kassimali §9.9).

Switch to demo code

What this demo shows

- The two systems represent the **same physics**, differing only in DOF ordering
- Dense solvers ignore sparsity and bandwidth entirely (unaffected)
- Larger bandwidth typically leads to higher factorization cost for sparse solvers
- Sparse solvers are sensitive to ordering, but can reduce its impact via reordering

Takeaway: DOF ordering changes the algebraic structure of \mathbf{K} , which can strongly influence computational cost even though the physics is unchanged.

Part 2 — Extending the DSM to 3D Trusses

Same DSM workflow, but with 3 translational DOFs per node and a 3D orientation.

Degrees of freedom in 3D

For a 3D truss node:

$$\mathbf{d}_i = [u_{xi} \quad u_{yi} \quad u_{zi}]^T, \quad \mathbf{d}_j = [u_{xj} \quad u_{yj} \quad u_{zj}]^T$$

Element displacement vector in global coordinates:

$$\mathbf{d}_e = [u_{xi} \quad u_{yi} \quad u_{zi} \quad u_{xj} \quad u_{yj} \quad u_{zj}]^T \equiv \text{DOFs } [1, 2, 3, 4, 5, 6]$$

3D direction cosines

Let the element connect nodes i and j with coordinates (x_i, y_i, z_i) and (x_j, y_j, z_j) .

$$L = \sqrt{(x_j - x_i)^2 + (y_j - y_i)^2 + (z_j - z_i)^2}$$

$$l = \frac{x_j - x_i}{L}, \quad m = \frac{y_j - y_i}{L}, \quad n = \frac{z_j - z_i}{L}$$

These are the direction cosines of the element axis in global coordinates.

3D Truss Element Matrices

We will write:

- the **local** element stiffness \mathbf{k}_{local} (6×6)
- the **transformation** matrix \mathbf{T} (6×6)
- the **global** element stiffness \mathbf{k}_{global} (6×6)

Local stiffness matrix (6×6)

A truss element carries **axial force only** and therefore has stiffness only in the direction of its axis.

In a local coordinate system where the element axis is the local x' direction, only the axial DOFs are active.

Then the 6×6 local stiffness is:

$$\mathbf{k}_{local} = \frac{EA}{L} \begin{bmatrix} 1 & 0 & 0 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Only axial coupling appears; transverse DOFs have zero stiffness in a truss model.

Transformation matrix \mathbf{T} (6×6)

For a truss, we only need the mapping from global translations to **axial** direction. A convenient 6×6 transformation (acting on translations only) is:

$$\mathbf{d}_{local} = \mathbf{T} \mathbf{d}_{global}$$

with

$$\mathbf{T} = \begin{bmatrix} \mathbf{Q} & \mathbf{0} \\ \mathbf{0} & \mathbf{Q} \end{bmatrix} = \begin{bmatrix} l & m & n & 0 & 0 & 0 \\ y_x & y_y & y_z & 0 & 0 & 0 \\ z_x & z_y & z_z & 0 & 0 & 0 \\ 0 & 0 & 0 & l & m & n \\ 0 & 0 & 0 & y_x & y_y & y_z \\ 0 & 0 & 0 & z_x & z_y & z_z \end{bmatrix}.$$

This version is sufficient because \mathbf{k}_{local} only acts on the axial rows/cols.

Note: A full 3D rotation basis would include two additional transverse unit vectors. For a truss, those transverse directions do not contribute stiffness.

Global element stiffness (6×6)

Compute:

$$\mathbf{k}_{global} = \mathbf{T}^T \mathbf{k}_{local} \mathbf{T}.$$

For a 3D truss element, the result can be written directly as:

$$\mathbf{k}_{global} = \frac{EA}{L} \begin{bmatrix} l^2 & lm & ln & -l^2 & -lm & -ln \\ lm & m^2 & mn & -lm & -m^2 & -mn \\ ln & mn & n^2 & -ln & -mn & -n^2 \\ -l^2 & -lm & -ln & l^2 & lm & ln \\ -lm & -m^2 & -mn & lm & m^2 & mn \\ -ln & -mn & -n^2 & ln & mn & n^2 \end{bmatrix}.$$

This is the standard 3D truss element stiffness in global coordinates.

Supports and stability in 3D

A free 3D structure has **3 rigid body translations**.

For a 3D **truss** model (translations only), you must prevent rigid body motion by constraining enough nodal translations.

Typical sufficient constraint patterns include (examples):

- Fix all three translations at one node, and two at a second node, and one at a third node (if geometry allows)
- Or use support conditions consistent with the physical supports (pins/rollers) but ensuring no global drift

If constraints are insufficient, **K** is singular → the model has a mechanism.